

A Novel RAG-Based Chatbot Solution to Improve Textbook Material Understanding

Vini Rezanejad
Computer Systems Lab
Thomas Jefferson High School for
Science and Technology
Alexandria, Virginia 22312
Email: radin.rezanezhad@gmail.com

Abstract—Mainstream large language models answer general questions from a user, but they aren't well-suited to semantically pick out relevant text from textbooks and lengthy PDF files, nor do they make use of this information retrieval towards proven active recall methods. This paper proposes a novel method that combines information retrieval through semantic search with active recall methods to create a chatbot that improves understanding of textbook content and the quality of study sessions. The implementation, which leverages a Retrieval-Augmented Generation (RAG) framework and a Rephrase-and-Respond (RaR) strategy, was evaluated across six AP-level subjects, achieving an average performance score of 0.906 out of 1.0. The chatbot demonstrated strong accuracy, reasoning, and memory retention, with particularly high performance in content areas such as History and Chemistry. These results highlight the potential of this system to provide credible academic assistance across a range of subjects and formats.

1. Introduction

Large Language Models (LLMs) have emerged as powerful tools in their ability to answer a wide array of questions spanning many subjects. In fact, a survey from Pew Research Center in February 2024 reveals that 43% of adults under the age of 30 have reported using ChatGPT in the past [5]. However, they are severely limited in the scope of education due to their inability to use and cite information from specific sources effectively. OpenAI's GPT-3.5, for example, when asked to cite sources for research purposes, fabricated or incorrectly cited sources in 55% of cases, revealing a significant risk of misinformation when models operate without grounded retrieval [8].

Additionally, they do not inherently incorporate successful learning strategies backed by research. These methods include the Feynman Technique, flashcards, practice questions, and other active recall strategies. In recent years, Retrieval-Augmented Generation (RAG) has gone beyond traditional LLMs in that they can access information from a database both accurately and quickly [4]. They serve as a great alternative to fine-tuning in the context of improving accuracy by referring to non-AI-generated information and

citing sources behind reasoning, marking a step forward for generative AI.

Currently, studies have already been using textbook information for the purpose of serving as a context behind generative chatbot retrieval. For instance, Singer et al. introduced a RAG-based chatbot model to leverage textbook knowledge in the field of ophthalmology using PDF copies of numerous ophthalmology textbooks, while also citing its textbook sources [7]. Alawwad et al. introduced a RAG-based approach specifically for textbook retrieval from middle school-level science-based textbooks, covering biology, earth science, and physics using a dataset containing small chunks of text and images, making way for up to a 9.84% improvement in model performance in multiple-choice questions [1]. Levonian et al. investigated middle school-level math and how RAG can generate better responses than traditional LLMs; however, the RAG didn't produce optimal results when questions were too grounded in the textbook [3].

One limitation common to all three of these studies is that textbook retrieval using RAG-based chatbots is limited to a single subject, and in some cases to a single textbook. The wide variety of textbooks used throughout different settings go beyond a single textbook, even among the same subject, making it more difficult for people to have a textbook-specific chatbot available to them. Additionally, these studies focus exclusively on answering questions from the textbook, even though there are much wider applications to using a chatbot based on a textbook, such as generation of flashcards and problem set generation. Furthermore, the benchmark performance for a RAG-based model is roughly 84% based on the results of these studies, suggesting that there is room for improvement on the performance of these models [1].

Recently, advancements have been made toward RAG-based models, one of which is getting the chatbot to rephrase questions before retrieving any information from the textbook, known as the Rephrase and Respond (RaR) method. When a question is phrased one way, there is a possibility that the LLM may not understand or interpret everything correctly, and answers may vary from the true answer. By rephrasing the question multiple times, we can improve output quality and significantly reduce GPT-3.5 hallucinations

by up 13.16% [2].

The purpose of this research study is to introduce an application that combines RAG with RaR strategies to improve reading retention and the quality of study sessions. To effectively prepare a method of study, the LLM needs to understand what content is being taught. The system will ask for the user to attach a file (or directory of files), allowing for full flexibility of what information the chatbot would like to retrieve.

2. Methods

All backend code for this project was written in Python. The frontend of the project, involving the chatbot GUI, was written in HTML, CSS, and JavaScript. A Python Flask app connects the frontend to the backend.




Figure 1. Diagram summarizing the custom RAG-based process based on user input

This implementation builds upon traditional RAG by also implementing RaR to make prompts clearer. Additionally, flashcard generation was implemented due to it being shown to increase student participation and academic performance [6]. It also allows the user to input any textbook file and supports back-and-forth chat to interact with the chatbot in a way that improves understanding. The custom RAG process is summarized in Fig. 1.

2.1. Text Extraction

The framework for the chatbot begins by extracting text from a source of information on a web server. On this server, the user is met with a home page asking the user to input a textbook file (Fig. 6). Here, this applies to data coming from a textbook either as a PDF, image, or audiobook file. The Flask server uses the POST method to take the PDF and allows all files to be added to the “uploads” folder in the server to enable server access.

2.1.1. PDF Files. The extracted text is split into chunks based on each page. PDFs used the PyMuPDF library on Python. PyMuPDF allows the code to extract text and keep track of the page while looping through the pages of a PDF file. The data structure used to keep track of the pages and text chunks is a dictionary. Fig. 2 displays the page-by-page algorithm of splitting text chunks, while Fig. 3 displays the use of dictionaries and its later conversion to CSV files. The `text_tokenized` variable contains all text in each page, and if this is not supported it uses image extraction methods with Pillow and pytesseract, explained in the next section. During this page-by-page loop, the page count goes up by one as each page number and text chunk is added to the dictionary.

```

book = pymupdf.open(filepath)
global totalpages
for i,page in enumerate(book):
    totalpages+=1
    text_tokenized = page.get_text().split()
    if not text_tokenized:
        pix = page.get_pixmap()
        image = Image.frombytes("RGB", [pix.width, pix.height], pix.samples)
        text_tokenized = pytesseract.image_to_string(image, lang='eng').split()
        if not text_tokenized: continue
    pages.append(totalpages)
    texts.append(" ".join(text_tokenized))
  
```

Figure 2. Algorithm for text extraction

```

dct = {"page":pages, "text":texts}
df = pd.DataFrame(dct)
  
```

Figure 3. Setting up data structure of parsed text and locations with a dictionary

2.1.2. Images. Much like PDF files, images are also split based on page. The key difference between image and PDF files is that images use the Pillow and pytesseract libraries to extract text that otherwise would be extracted from a PDF file. To enable input of multiple images or PDF files, all the image files are converted to PDF files using Pillow to enable better compatibility using the code snippet shown in Fig. 4. The outputted PDF file is also contained in the “uploads” folder. The server supports use of multiple PDF or image files. If a PDF file also does not support the extraction of text, it also uses Pillow and pytesseract to extract all text in case pymupdf is unable to parse PDF text.

```

def img_to_pdf(image_paths, out_pdf):
    imgs = [Image.open(p).convert("RGB") for p in image_paths]
    imgs[0].save(out_pdf, save_all=True, append_images=imgs[1:])
  
```

Figure 4. Code snippet showing the image-to-PDF conversion process

2.1.3. Audiobooks. The extracted text is split based on 30-second paragraphs from an MP3 file using the Deepgram API. Deepgram is capable of extracting text from an audiobook file based on paragraphs and sentences. An example

of sentences in the JSON object can be shown in Fig. 5, where paragraphs have similar functionality but are much longer. Extraction was chosen to be based off paragraphs, because sentences tended to be too short to be useful both practically and semantically. Deepgram's JSON object also contains the timestamps of the paragraphs including start and end times, allowing us to use the start timestamp for the user's reference. Since these times were given in seconds, it was necessary to convert in terms of hours, minutes, and seconds for more practical use from the user's point of view. The server supports use of one audiobook file.

```

    "end": 3178.975,
    "num_words": 188
  },
  {
    "sentences": [
      {
        "text": "This much I understand now.",
        "start": 3173.551,
        "end": 3173.7751
      },
      {
        "text": "But the man can give no help to the boy, not in this matter nor in those that follow",
        "start": 3181.95,
        "end": 3181.95
      },
      {
        "text": "The boy moves always out of reach.",
        "start": 3182.3699,
        "end": 3182.3699
      },
      {
        "text": "One afternoon, I walked a friend of mine to his house.",
        "start": 3181.655,
        "end": 3181.655
      },
      {
        "text": "After he went inside, I sat on his steps for a while, then got to my feet and started toward home, we",
        "start": 3192.535,
        "end": 3201.55
      }
    ],
    "start": 3171.7551,
    "end": 3201.55,
    "num_words": 65
  },
  {
    "sentences": [

```

Figure 5. Example of Deepgram text JSON, with start/end times for sentences

When the user submits the file(s), the server will load, going through extraction and embedding conversion, before going to the chat page. This takes roughly 15 seconds for a 1000-page PDF file, or 20 seconds for a one-hour audio file. This data, storing both the text and the location of the text, is stored into a CSV file, called “texts.csv”, by concatenating the text with either the page number or the timestamp. If this CSV file already exists, it is deleted and recreated to make way for user-inputted textbook information even after it was used in a previous session.




Figure 6. Home page of the web server, where the user inputs textbook file(s)

2.2. Embeddings

After storing the data into a CSV file, each text chunk from the textbook is converted to a 384-dimensional vector

embedding. Using pandas once again this time for retrieval of CSV data, each chunk of text is tokenized into its words, inputted into this self-attention model, and outputs an embedding vector, with its magnitude and direction depending on the semantic definition of the words into the text. Similar text will be close together in this vector space and would be reflected after encoding all text information from the CSV file. In this specific use case, a 6-layer pre-trained Sentence-BERT self-attention model was used, generating a vector space based on semantics. Upon completion of this step, the chat page loads (Fig. 7), ready for the user to input a question.

2.3. RaR Implementation

When the user inputs a question, the chatbot rephrases the question into one that is clearer and more detailed using the OpenAI GPT-3.5 model, leaving it less prone to hallucinations by essentially asking itself the question. Still, even after the question is rephrased, the original question remains as part of the prompt because there is a risk of the rephrased question taking away information from what the original question aimed to ask. This is prompted to GPT-3.5 as a “user” prompt, as it is the user aiming to get the question rephrased. The RaR prompt is shown in Fig. 8.




Figure 7. Visual of the chat page after the server processes text chunks and embeddings

```
{  
    "role": "user",  
    "content": f'Rephrase the following question to be more detailed and clearer:\n\n{question}'  
}
```

Figure 8. Code snippet displaying GPT-3.5 prompt used to implement RaR

2.4. Retrieval

This rephrased question is also converted to an embedding of its own using the same SBERT embedding model and is compared to the 5 nearest embedding vectors in the text using the cosine similarity metric. This is done by taking all the cosine similarity values and sorting it alongside the respective embeddings. Those 5 nearest embeddings will then be selected as retrieved text for GPT-3.5's prompt. The

reason behind using GPT-3.5 is not only because it is one of the most cost-effective large language models commonly used, but it is also a well-researched model that can be clearly compared with existing literature in evaluating the chatbot.

2.5. Prompting and Output

As this data is extracted, it is also part of the GPT-3.5 prompt to answer the user's question. Included in the prompt includes the following data from the user: original question, rephrased question, retrieved text with its sources cited by the page number or timestamp. Additional prompting information includes prompting the chatbot to use step-by-step reasoning, simple language, clear and detailed explanations, wrapping LaTeX expressions with “ $\$\$$ ” for display equations or “ $\$$ ” for inline math. This is done as part of the “system” role in GPT’s API as it is meant to set the behavior of GPT-3.5 when the user asks the question.

After the server loads the GPT-3.5 generated response, the user interface will display both the user's question and the chatbot's responses are formatted in a readable format, especially with expressions in subjects including but not limited to Math and Chemistry, the MathJax library is used in JavaScript to convert LaTeX code to readable text. The reason behind using the dollar signs in the previous step was for MathJax to identify LaTeX expressions so that it can correctly process any equations, expressions, or any special characters requiring LaTeX format. The result on the HTML page can be shown with the example of the steps of u-substitution on Figure 9. This is part of the “assistant” role as it is indeed the assistant providing the generated response. At the same time, from the frontend, a series of JavaScript functions are run so that the chat box, treated as a list of chat elements, keeps adding on list items (specifically, back-and-forth interactions between the user and the chatbot) and styles them differently depending on the role contained in each message of the chatbot's history stored in the cookie. Additionally, fetch was used on the Flask server to ensure that the user stays on the same page while the chatbot information is being processed.

The screenshot shows a browser window with a tab labeled "Textbooks...". The main content area displays a step-by-step guide for solving the integral $\int x^2 \sqrt{1-x^2} dx$ using u-substitution. Step 1: Looks for a nested composition $f(g(x))$ within the integrand for which the inner function is a power of the outer function. Step 2: Looks for a common factor in the numerator and denominator. Step 3: If you are successful for Step 2, then replace $u = g(x)$ to express your final answer in terms of u . The guide then provides a detailed example of the substitution process, showing the transformation of the integral into terms of u and du , and the final evaluation of the integral.

Figure 9. Chatbot output using MathJax with citations after asking the chatbot to explain u-substitution

2.6. Flashcard Generation

Beyond the RAG framework, this chatbot also utilizes generative flashcards, enhancing recall in the user by quizzing themselves on a topic they choose. This is done by retrieving related information to the inputted topic and generating questions and answers. Fig. 10 shows an example of 10 flashcards being generated on the topic of derivatives.

The screenshot shows a web application titled "Generate Flashcards" with the topic "Derivatives". It displays 10 flashcards arranged in a grid. Each card has a question (Q) on the front and an answer (A) on the back. The cards cover various topics related to derivatives, such as the derivative of sine, chain rule applications, and higher-order derivatives. A "Create" button is visible in the top right corner.

Figure 10. Demonstration of flashcard generation, with front (Q) and back (A) sides displayed upon clicking them

2.7. Memory

Back-and-forth memory was implemented, allowing the chatbot to accommodate any special requests made by the user, including but not limited to response length, language complexity, and analogies. This was done by creating a cookie in the Flask server which contains all history information of the back-and-forth interaction between chatbot and user, including the system instructions so the chatbot remembers what it is meant to do. Specifically, this cookie is a dictionary with a “history” attribute which contains the entire history of questions and answers. Modifying this is displayed in Fig. 11. When a user inputs a new question, the new question and response are simply appended to the existing memory and is stored again into the cookies once a response is generated. Fig. 12 demonstrates an example of back-and-forth chat working with an AP Calculus BC textbook, after asking the chatbot to summarize its original output in two sentences.

```
history.append({"role": "user", "content": question})
history.append({"role": "assistant", "content": answer})
session["history"] = history
```

Figure 11. Code snippet modifying the Flask server cookie to add chatbot memory

2.8. Evaluation Criteria

The model was evaluated qualitatively using a combination of metrics that assess its performance in multiple criteria. The criteria with respective weighting in parentheses based on importance are:

The figure shows a back-and-forth conversation between a user and a chatbot. The user's message is: "How do I use u-substitution?". The chatbot's response is a detailed step-by-step guide on how to use u-substitution, including formulas and examples. The user then asks: "Explain this in two sentences." The chatbot responds with a summary of the concept and its application.

Figure 12. Demonstration of back-and-forth chat, asking the chatbot to follow-up its original response with a brief summary

- 1) Is the information accurate? (0.3)
- 2) Does it account for complexities in the facts? (0.05)
- 3) Does it cite its sources? (0.1)
- 4) If it cites, does the source reflect what is being outputted? (0.15)
- 5) Is it easy to understand? (0.05)
- 6) Does it provide the right reasoning? (0.1)
- 7) Does it explain applications well? (0.05)
- 8) Is it accommodating to special requests by the user? (0.1)
- 9) Does it recall past chats well enough? (0.1)

The model was evaluated on 6 different subjects: AP Chemistry, AP US History, AP Physics C, AP Calculus BC, AP Computer Science, and AP Government. All these subjects were tested at the AP level, and for AP Government an audiobook was used to test the chatbot’s performance in a different format. For each subject, the chatbot was asked 10 questions, each in a different conversation. After this question, a relevant follow up question was asked, which was used to evaluate the model’s conversational ability.

3. Results and Discussion

TABLE 1. EVALUATION SCORES ACROSS SUBJECTS (OUT OF CATEGORY WEIGHTS)

Topic	Acc.	Comp.	Cit.	Corr.	Easy	Reas.	App.	Sp.	Mem.	Tot.
Chemistry	0.28	0.05	0.098	0.14	0.041	0.10	0.012	0.10	0.10	0.921
History	0.30	0.05	0.09	0.125	0.047	0.095	0.05	0.095	0.10	0.952
Physics	0.30	0.05	0.07	0.105	0.05	0.10	0.05	0.10	0.10	0.895
Math	0.30	0.05	0.07	0.095	0.05	0.095	0.032	0.095	0.09	0.877
Comp Sci	0.29	0.05	0.07	0.082	0.046	0.10	0.05	0.10	0.10	0.888
Gov Audio	0.28	0.04	0.10	0.138	0.05	0.10	0.045	0.08	0.07	0.903
Avg.	0.292	0.048	0.083	0.114	0.047	0.098	0.040	0.095	0.093	0.906

The results of the chatbot’s performance can be summarized in Table I. The chatbot achieved strong overall performance with an average normalized score of 0.906 across all subjects. The highest-scoring subjects were U.S. History (0.952 normalized), Chemistry (0.921 normalized), and Government (0.903 normalized), indicating that the chatbot excels in subjects where discrete factual recall and

empirical reasoning dominate. These subjects typically feature less repetitive scaffolding between concepts, allowing the RAG framework to extract and cite localized information more effectively.

In contrast, performance dipped slightly in Math (0.877 normalized) and Computer Science (0.888 normalized). These subjects often involve more interdependent knowledge structures, i.e. where theorems build on prior proofs or code requires multi-line context. This complexity posed challenges for accurate citation and chunk-based semantic retrieval. Specifically, CS suffered from a noticeable drop in correct citation scores, which can be attributed to GPT-3.5’s default bias toward code synthesis rather than source-grounded explanations.

The overall citation inclusion scored 0.083/0.1, while correct citation accuracy reached 0.114/0.15. These are substantial improvements over traditional GPT-3.5 performance. As mentioned previously, GPT-3.5 fabricated citations in over 50% of cases [8]; in contrast, the RAG-based method achieved a citation correctness score of 0.76 (normalized by weight) when citations were present. What was especially impressive about the result, though, is that there was a citation presence score of 0.838 when normalized, demonstrating that there is a higher likelihood of correct citations relative to any citation being present.

The Government audiobook case revealed unique trends. Despite using less granular chunking (30-second intervals), it maintained a high total score. This may be from a smaller semantic space: because the source content is more limited, retrieved passages are semantically closer to queries by default. However, this advantage came at the expense of memory performance, particularly for longer conversations where timestamped context became less precise.

4. Conclusion

This chatbot demonstrates strong performance when it comes to meeting the 9 criteria, especially when it comes to generating accurate information that is easy to understand and provides good reasoning. It also proved to be strong in citing its sources, a key feature. However, there were some limitations, especially when it came to retrieving the right information in audiobooks, as well as citing the right sources in material that involved code, formula generation, and interdependent knowledge.

In the future, studies could involve more optimized methods of chunk splitting from the original source, especially with textbook PDFs, that split based on topic rather than by page, to differentiate the various embeddings as much as possible. Additionally, future studies should use newer and more cost-efficient large language models than the models used here.

Acknowledgments

The author would like to thank Dr. Torbert for his guidance through the TJHSST Computer Systems Lab.

References

- [1] H. A. Alawwad, A. Alhothali, U. Naseem, A. Alkhathlan, and A. Jamal, "Enhancing textual textbook question answering with large language models and retrieval augmented generation," *Pattern Recognition*, vol. 162, p. 111332, 2025.
- [2] Y. Deng, W. Zhang, Z. Chen, and Q. Gu, "Rephrase and respond: Let large language models ask better questions for themselves," *arXiv preprint arXiv:2311.04205*, 2023.
- [3] Z. Levonian, C. Li, W. Zhu, A. Gade, O. Henkel, M. E. Postle, and W. Xing, "Retrieval-augmented generation to improve math question-answering: Trade-offs between groundedness and human preference," *arXiv preprint arXiv:2310.03184*, 2023.
- [4] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, et al., "Retrieval-augmented generation for knowledge-intensive NLP tasks," in *Advances in Neural Information Processing Systems*, vol. 33, pp. 9459–9474, 2020.
- [5] C. McClain, "Americans' use of ChatGPT is ticking up, but few trust its election information," Pew Research Center, March 2024.
- [6] Y. S. Nugroho, J. Nurkamto, and H. Sulistyowati, "Improving students' vocabulary mastery using flashcards," *English Education*, vol. 1, no. 1, 2012.
- [7] M. B. Singer, J. J. Fu, J. Chow, and C. C. Teng, "Development and evaluation of Aeyeconsult: A novel ophthalmology chatbot leveraging verified textbook knowledge and GPT-4," *Journal of Surgical Education*, 2023.
- [8] W. H. Walters and E. I. Wilder, "Fabrication and errors in the bibliographic citations generated by ChatGPT," *Scientific Reports*, vol. 13, no. 1, 2023.